

EXERCISES

TASK #1

Problem statement:

We need to connect to the SQL Server and create a bank account database using the following SQL script.

Solution:

Copying the following code :

Code:

```
Use lab4
create table account
(id int primary key,
balance int)
insert into account values(1, 5000)
insert into account values(2, 8500)
insert into account values(4, 6400)
```

Reasoning: Result:

We have got an account table with different balances for each account.

TASK #2

Problem statement:

We need to create a transaction to transfer 1000 HUF from bank account n°1 to n°2.
In order to check the transfer open another query window, this creates a second parallel connection to the database.

Solution:

By Data Definition Language scripts :

We use UPDATE command to modify the existing records in the account table.

And in order to specify which columns and values that should be updated in a table we use SET command :

* We subtract from the balance of the 1st account by 1000huf.

* Then we add 1000huf to the balance of the 2nd account.

Code:

```
(1)  select * from account

(2)  use lab4
      update account
      set balance=balance-1000
      where id=1

(3)  select * from account

(4)  use lab4
      update account
      set balance=balance+1000
      where id=2

(5)  select * from account
```

Reasoning: Result:

	id	balance
1	1	5000
2	2	8500
3	4	6400

(1)

Messages

(1 ligne affectée)

Heure de fin : 2022-04-11T23:55:14.3432439+02:00

(2)

	id	balance
1	1	4000
2	2	8500
3	4	6400

(3)

Messages

(1 ligne affectée)

Heure de fin : 2022-04-11T23:56:39.3604101+02:00

(4)

	id	balance
1	1	4000
2	2	9500
3	4	6400

(5)

We have got an updated table with a new balance for each operation.

We notice that a 1000huf was transferred from the 1st row (account 1) to the 2nd row (account 2)

TASK #3

Problem statement:

We need to do again the second task by changing the script in a way that the data manipulation operations are put into the same transaction.

Solution:

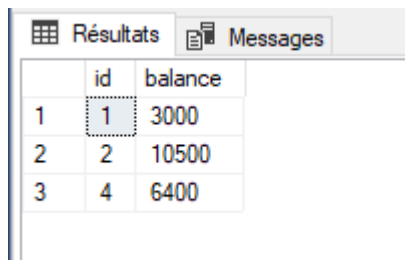
By Data Definition Language scripts :

Code:

```
Use lab4
begin transaction
update account
set balance=balance-1000 where id=1
update account
set balance=balance+1000 where id=2
```

Reasoning:

Result:



	id	balance
1	1	3000
2	2	10500
3	4	6400

We notice that we got a different result compared to the previous task, that is because the transaction has not yet been committed while the system were trying to read the data. That is called Dirty read.

TASK #4

Problem statement:

We need to do two transfers in parallel: one of them sends 500 HUF from account 1 to 2 while the second one sends 300 HUF from bank account 2 to 1.

Solution:

By Data Definition Language scripts :
We copy the following code

Code:

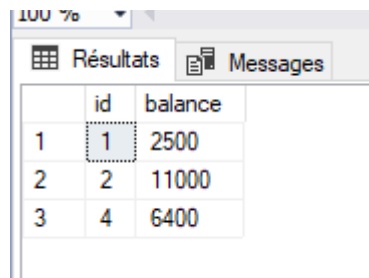
```
use lab4
begin transaction
update account
set balance=balance-500
where id=1
update account
set balance=balance+500
where id=2
```

First transaction

```
use lab4
begin transaction
update account
set balance=balance-300
where id=1
update account
set balance=balance+300
where id=2
```

Second transaction

Reasoning: Result:



	id	balance
1	1	2500
2	2	11000
3	4	6400

Each transaction cannot continue and begins to wait for the other one, that is why we got the same result. That is called deadlock.

TASK #5

Problem statement:

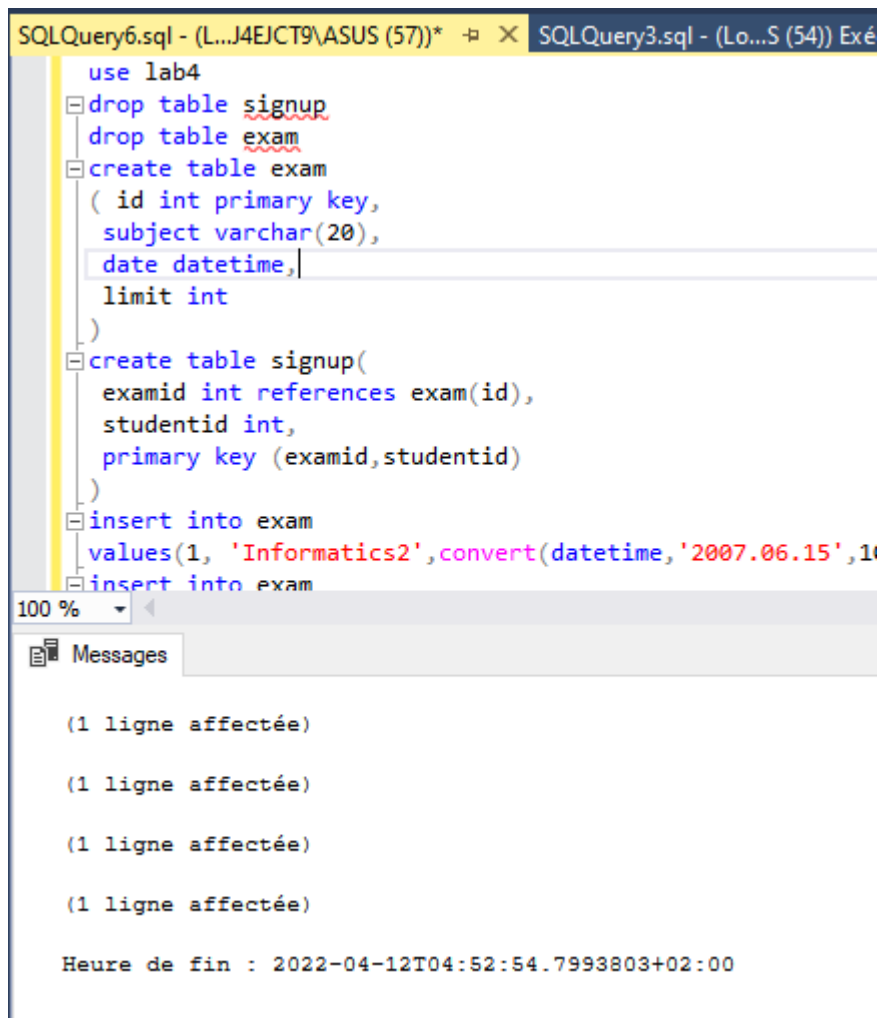
We need to create a simple exam signup system using the following SQL script:

Solution:

By Data Definition Language scripts :

We need to copy the following code

Code and Result:



The screenshot shows a SQL Server Enterprise Manager window with two tabs: 'SQLQuery6.sql - (L...J4EJCT9\ASUS (57))*' and 'SQLQuery3.sql - (Lo...S (54)) Exé'. The active tab displays a SQL script with the following content:

```
use lab4
drop table signup
drop table exam
create table exam
( id int primary key,
  subject varchar(20),
  date datetime,
  limit int
)
create table signup(
  examid int references exam(id),
  studentid int,
  primary key (examid,studentid)
)
insert into exam
values(1, 'Informatics2',convert(datetime,'2007.06.15',1)
insert into exam
```

Below the script, the 'Messages' pane shows the execution results:

```
(1 ligne affectée)
(1 ligne affectée)
(1 ligne affectée)
(1 ligne affectée)
Heure de fin : 2022-04-12T04:52:54.7993803+02:00
```

Reasoning:

We have got table of an exam signup system.

TASK #6

Problem statement:

We need to simulate two concurrent signups to the first exam.

Solution:

By Data Definition Language scripts :

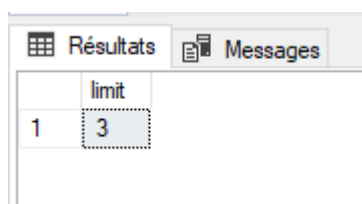
We need to copy the following code

Code:

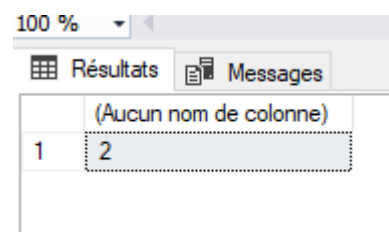
<pre>(1) use lab4 select limit from exam where id=1 (2) use lab4 select count(*) from signup where examid=1 (3) use lab4 insert into signup values(1, 333)</pre>	<pre>(1) use lab4 select limit from exam where id=1 (2) use lab4 select count(*) from signup where examid=1 (3) use lab4 insert into signup values(1, 444)</pre>
--	--

First signup

Second signup

Reasoning:**Result:**


	limit
1	3



	(Aucun nom de colonne)
1	2

We have got an error ,since we were repeating a read operation on the same records, but we got different records in the results set. This is called phantom read

TASK #7

Problem statement:

We need to do again the previous task by a different method

Solution:

We need to copy the new code and increase isolation levels of transactions to "serializable"

Code:

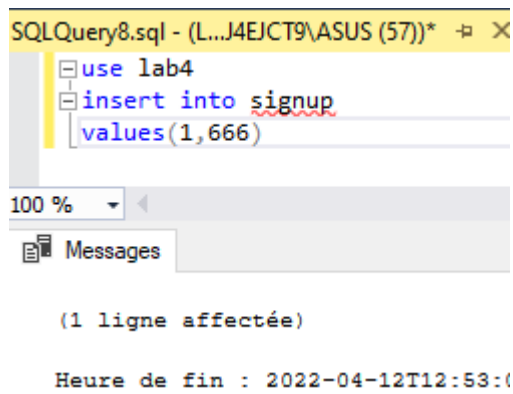
<pre>(1) set transaction isolation level serializable Begin transaction (2) select limit from exam where id=1 (3) select count(*) from signup where examid=1 (4) insert into signup values(1,555)</pre>	<pre>(1) set transaction isolation level serializable Begin transaction (2) select limit from exam where id=1 (3) select count(*) from signup where examid=1 (4) insert into signup values(1,666)</pre>
1 st signup	2 nd signup

Reasoning:

Result:

Résultats	Messages
limit	(Aucun nom de colonne)
1 3	1 4

Laboratory Report – Informatics 2



The screenshot shows a SQL query window titled 'SQLQuery8.sql - (L...J4EJCT9\ASUS (57))*'. The query text is: `use lab4`, `insert into signup`, and `values(1,666)`. Below the query, a status bar indicates '(1 ligne affectée)' and 'Heure de fin : 2022-04-12T12:53:00'. A 'Messages' tab is visible at the bottom.

The written values were added successfully

As we used serializable level, both of the transactions behave as if they were executed one after the other

TASK #8

Problem statement:

We need to increase limit of the first exam to 5, then do again the task7.

Solution:

Code: increase limit of the first exam to 5

```
use lab4
update exam set limit=5 where id=1
```

Reasoning:

Result:

Résultats	Messages				
<table> <tr> <th>limit</th></tr> <tr> <td>15</td></tr> </table>	limit	15	<table> <tr> <th>(Aucun nom de colonne)</th></tr> <tr> <td>16</td></tr> </table>	(Aucun nom de colonne)	16
limit					
15					
(Aucun nom de colonne)					
16					

We have got an increased limit and count.

```
Messages
Msg 2627, Niveau 14, État 1, Ligne 2
Violation of PRIMARY KEY constraint 'PK__signup__F1BD43411759BB85'.
The statement has been terminated.
```

We have got an error because we can not insert the same value multiple times.

That is why we have to change the inserted value to (2,555)and(2,666) because we already used (1,555) and(1,666) because we can not insert the same value multiple times.

```
SQLQuery2.sql - (L...J4EJCT9\ASUS (56))*
use lab4
insert into signup
values(2,555)

SQLQuery2.sql - (L...J4EJCT9\ASUS (5
use lab4
insert into signup
values(2,666)

100 %
Messages
(1 ligne affectée)
```

TASK #9

Problem statement:

We need to do again the task 7 with a different method :the first student should signup to the first exam while the second one should signup to the second exam.

Solution:

We need to recopy the code from the task 7

Code:

```
use lab4
```

Reasoning:

Result:

limit
1 5

limit
1 3

1st student:limit

2nd student: limit

(Aucun nom de colonne)
1 5

(Aucun nom de colonne)
1 0

1st student : count

2nd student : count

We have to change the inserted value to (3,555)and(3,666) because we already used (1,555) ,(1,666), (2,555)and(2,666) because we can not insert the same value multiple times.

Otherwise we get a result as error like the previous task

TASK #10

Problem statement:

We need to reset our database by executing the script of example 5 again.

Solution:

We copy and execute again the script from the task 5.

Code:

```
use lab4
drop table signup
drop table exam
create table exam
( id int primary key,
  subject varchar(20),
  date datetime,
  limit int
)
create table signup(
  examid int references exam(id),
  studentid int,
  primary key (examid,studentid)
)
insert into exam
values(1, 'Informatics2',convert(datetime,'2007.06.15',102),3)
insert into exam
values(2, 'Mathematics',convert(datetime,'2007.06.18',102),3)
insert into signup values(1,111)
insert into signup values(1,222)
```

Reasoning:

Result:

We recopy the script from task 5.

TASK #11

Problem statement:

We need to increase by using read committed isolation level and mutual locking.
If concurrent transactions lock only the record they need, mutual exclusion can be achieved.

Solution:

We copy the following code :

<pre> (1) set transaction isolation level read committed begin transaction (2) select limit from exam with(XLOCK) where id=1 (3) select count(*) from signup where examid=1 (4) insert into signup values(1,333) (5) commit </pre>	<pre> (1) set transaction isolation level read committed begin transaction (2) select limit from exam with(XLOCK) where id=1 (3) select count(*) from signup where examid=1 (4) commit </pre>
First signup	Second signup

Reasoning:**Result:**

	limit
1	5

	(Aucun nom de colonne)
1	3

The transaction holds a read on the current row, and prevents other transactions from reading it.

Since the isolation level guarantees to read only committed data, it does not allow dirty read.

TASK #12

Problem statement:

We need to increase limit of the first exam to 6, then do again the task11.

Solution:

Code: increase limit of the first exam to 6.

```
use lab4
update exam set limit=6 where id=1
```

Recopy the code from the previous task.

Reasoning:

Result:

Résultats		Messages	
limit			
1	6		

get an increased limit for the exams

Résultats		Messages	
(Aucun nom de colonne)			
1	2		

Count from 1st signup

Résultats		Messages	
(Aucun nom de colonne)			
1	3		

Count from 2nd signup

We have to change the inserted value to (2,333) because we already used (1,333) and we can not insert the same value multiple times.

TASK #13

Problem statement:

We need to do again the task 12 with a different method : the first student should signup to the first exam while the second one should signup to the second exam

Solution:

We copy the following code :

<pre>(6) set transaction isolation level read committed begin transaction (7) select limit from exam with(XLOCK) where id=1 (8) select count(*) from signup where examid=1 (9) insert into signup values(1,333) (10) commit</pre>	<pre>(5) set transaction isolation level read committed begin transaction (6) select limit from exam with(XLOCK) where id=2 (7) select count(*) from signup where examid=2 (8) commit</pre>
First signup	Second signup

Reasoning:

Result:

The second student should signup to the second exam => change the id and the examid from 1 to 2.

We also have to change the inserted value to (3,333) because we already used (1,333) and(2,333) and we can not insert the same value multiple times.